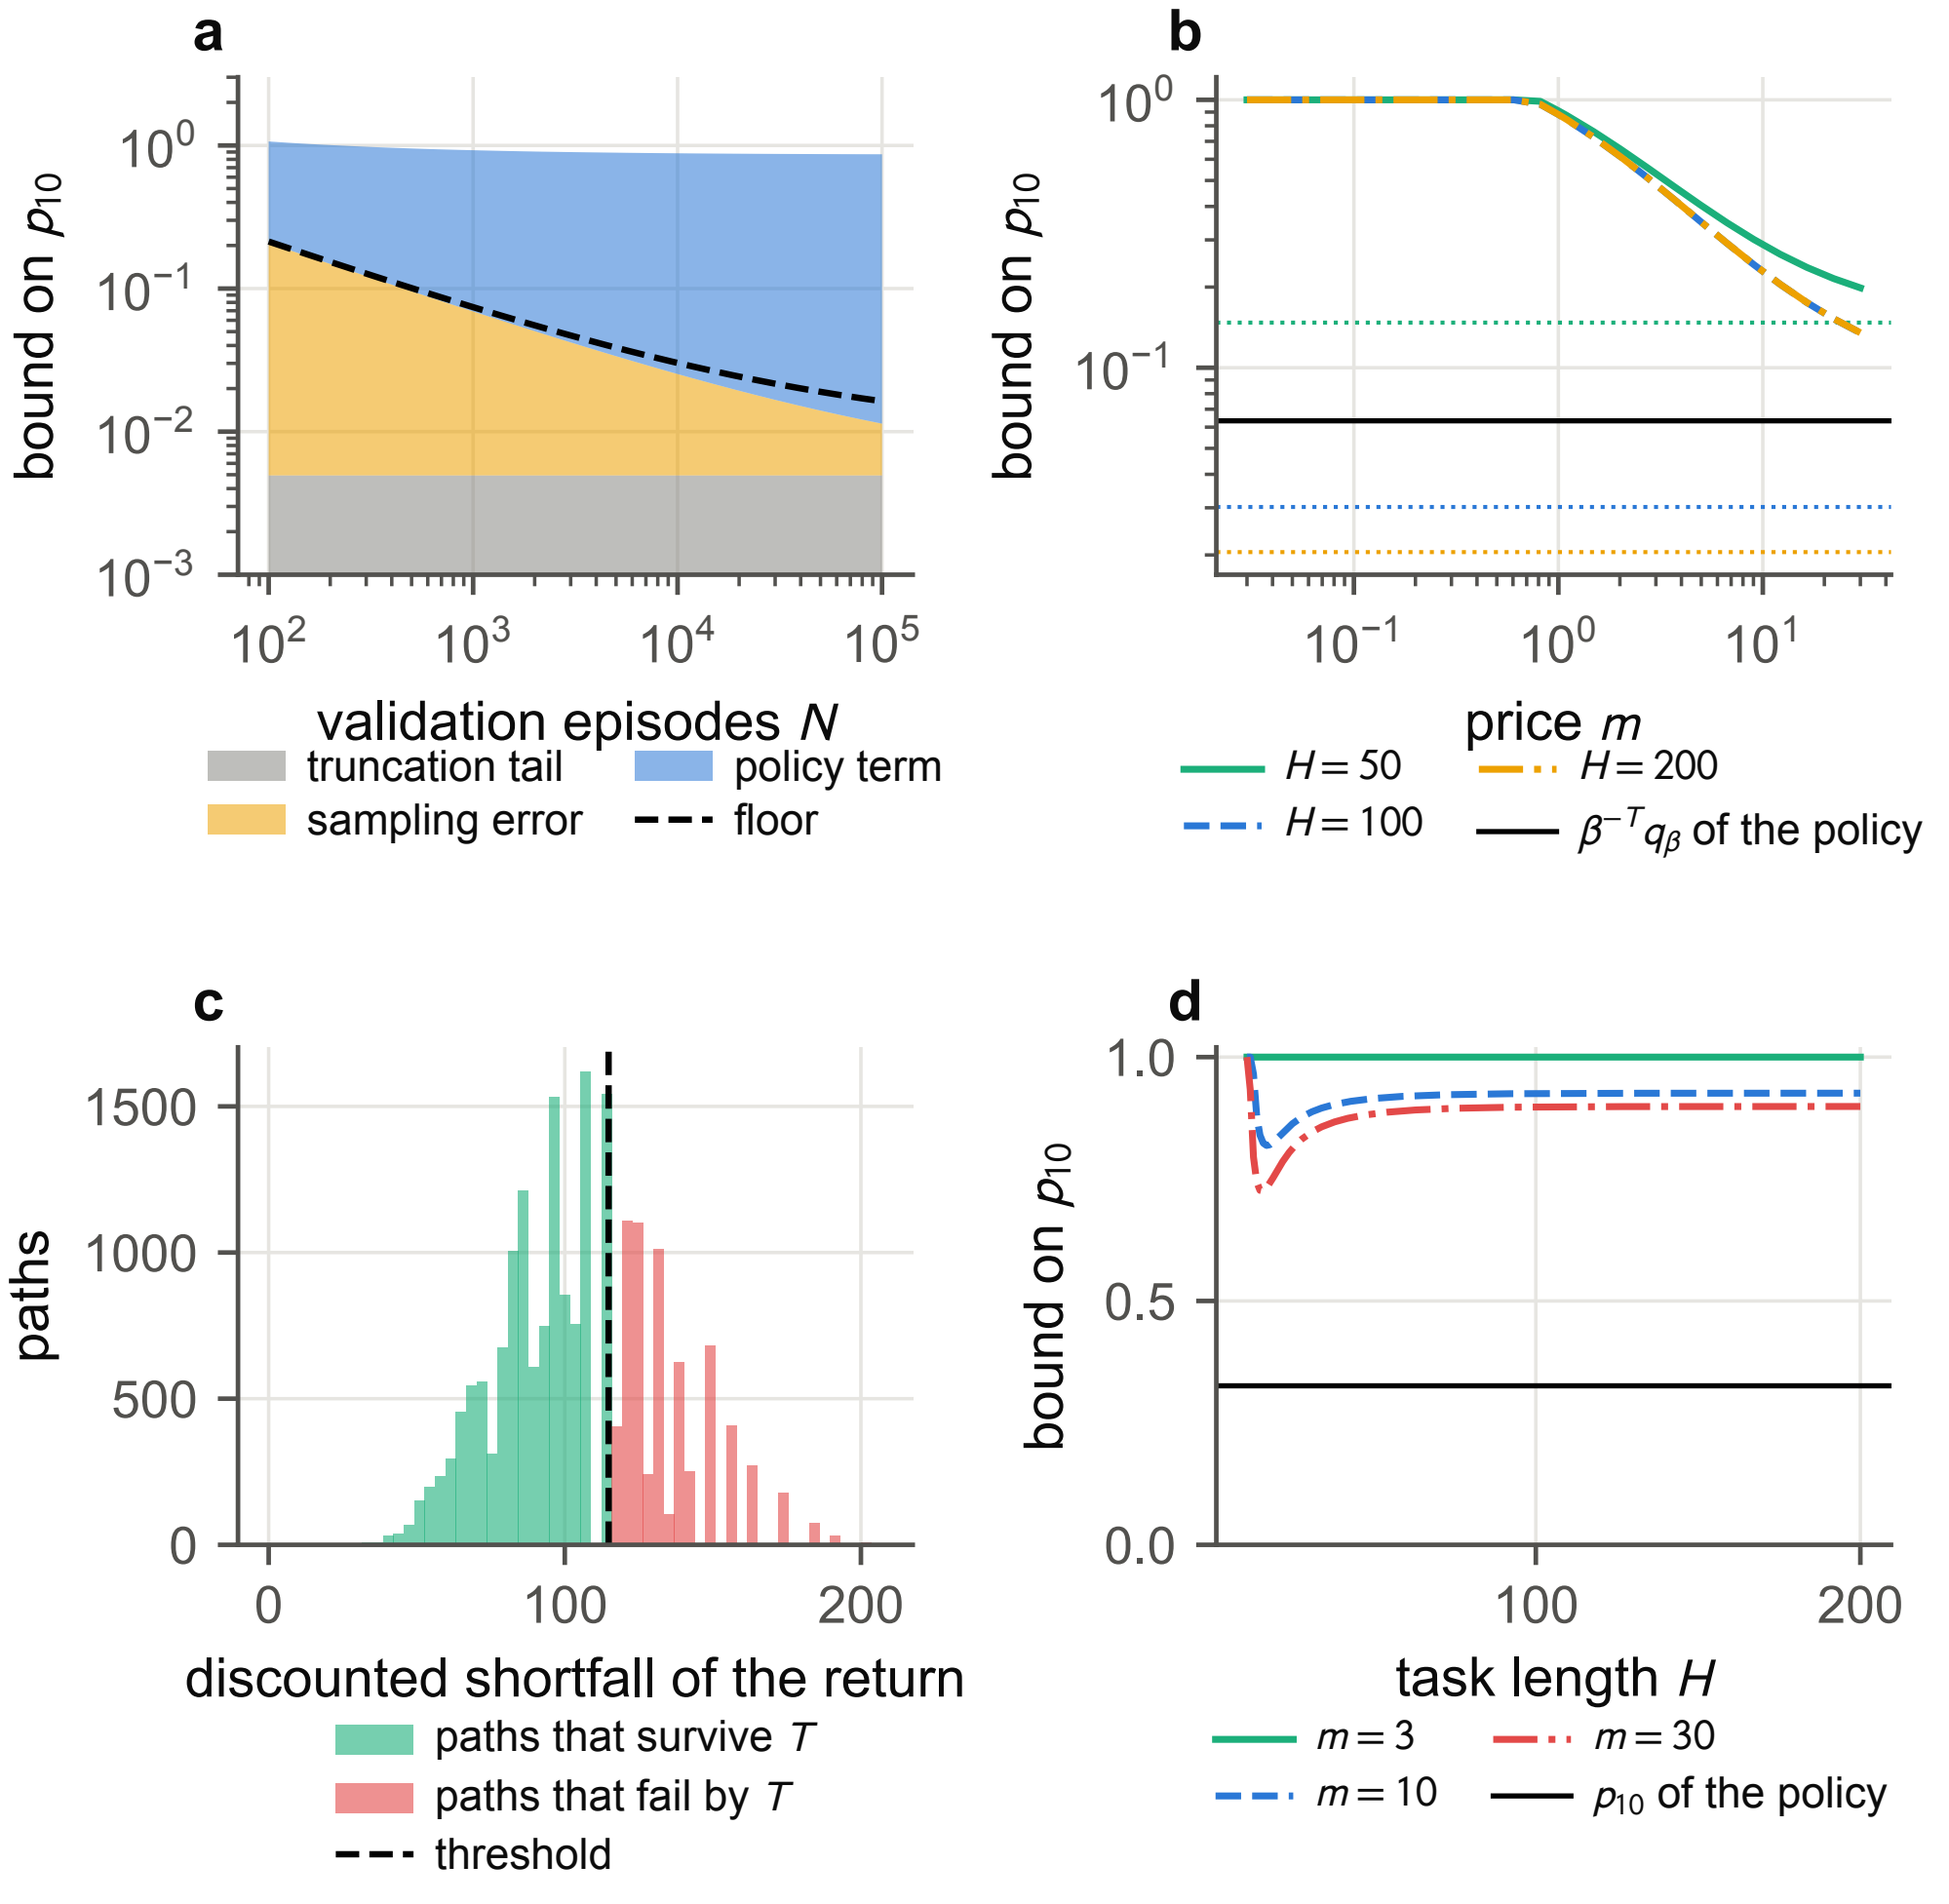
a
bound on p_{10}
10^0
10^{-1}
10^{-2}
10^{-3}
10^2
10^3
10^4
10^5
validation episodes N
truncation tail
policy term
sampling error
floor
b
bound on p_{10}
10^0
10^{-1}
10^{-1}
10^0
10^1
price m
$H = 50$
$H = 200$
$H = 100$
$\beta^{-T} q_\beta$ of the policy
c
paths
1500
1000
500
0
0
100
200
discounted shortfall of the return
paths that survive T
paths that fail by T
threshold
d
bound on p_{10}
1.0
0.5
0.0
100
200
task length H
$m = 3$
$m = 30$
$m = 10$
p_{10} of the policy